Performing the Nation as Commodity: Cyber-Nationalist Witch-Hunts around Wuchang: Fallen Feathers in China's Platformized Media Ecology

This study re-theorizes nationalism in the information age by examining how cyber-nationalist discourse fuses with platform economic logics to produce repeatable, commodified "witch-hunt" repertoires. Using the 2025 release of the single-player game Wuchang: Fallen Feathers (明末) as a key case, the project analyzes how accusations, moral panic, and intra-national boundary policing around the game were staged, circulated, and monetized across Chinese digital platforms. We argue that contemporary cyber-nationalism in China is not only an ideological project but also a patterned content-production logic: nationalist discourse becomes a template for packaging "witch-hunt" performances into scalable, cross-platform commodities.

We conceptualize nationalism as a form of identity discourse and emotional capital organized around the construction and policing of "us/them" boundaries. In an information environment structured by attention scarcity and algorithmic distribution, platform logics simultaneously promote polarization (the sharpening of antagonistic positions) and pluralization (the proliferation of competing narratives). At the same time, discursive authority has partially shifted from state–legacy media hierarchies to multi-nodal information flows that include short-video feeds, bullet chats, and community forums. Within this expanded arena, othering mechanisms—intensified by metrics and recommendation systems—amplify aggressive and accusatory tendencies, facilitating cyber "witch-hunts" structured around exposure, encirclement, and moral judgment.

Han nationalism is central to this configuration. Its identity core activates online role performance: users can strategically inhabit positions such as whistle-blower, defender of the nation, or moral judge. By performing these roles, individuals accumulate moral capital (appearing righteous, vigilant, and patriotic) while benefiting from platform rewards such as visibility, engagement, and potential monetization. Nationalism thus supplies both the symbolic vocabulary and the affective energy that make witch-hunt performances legible, emotionally compelling, and algorithmically attractive.

Wuchang: Fallen Feathers provides a theoretically rich site to observe these dynamics. Because the game engages late-Ming history and intersecting ethnic identities, it became a focal point for disputes over "proper" national representation and historical dignity. Across Bilibili, Douyin, Xiaoheihe, Zhihu, and related platforms, we observe accusations structured around Han nationalism, ethnic purity, and betrayal of the nation. The incident generates routinized content patterns: creators produce highly similar videos and posts, recycle talking points, and cross-post materials to capture attention in multiple arenas. These patterns suggest the emergence of a replicable content industry chain built around an "othering affect – antagonistic narrative – traffic monetization" pipeline.

On this basis, we pose three research questions. RQ1 asks whether nationalist witch-hunts follow a visually traceable, structurally stable, and replicable content logic, and what their key narrative units and transition rules are. RQ2 asks how this content logic is packaged, circulated, and consumed across platforms—specifically, how emotions are converted into tradeable "traffic-to-revenue" commodities through platform mechanisms. RQ3 probes the individual motivations that underpin participation in witch-hunt performances and their associated content industry chain.

Empirically, we construct a cross-platform corpus of *Wuchang: Fallen Feathers*–related information, reviews, commentary videos, discussion threads, comments, and bullet chats from Bilibili, Douyin, Xiaoheihe, and Zhihu. The collection window spans from launch week through the decline of public attention, capturing the full lifecycle from initial trigger to saturation and cooling. The unit of analysis includes both content items (videos, posts, threads) and interactional traces (comments, replies, bullet chats), allowing us to track how narratives circulate and mutate across platforms.

Methodologically, the project combines computational text analysis with theory-driven discourse coding. First, we use Chinese pre-trained language models (RoBERTa) and sentence-level embeddings (SBERT) to compute semantic similarity and cluster texts into recurring narrative and evaluative patterns. Second, we develop a lexicon and labeling scheme for nationalist frames—such as national dignity, Han nationalism, ethnic purity, betrayal, external and internal enemies—and implement a hybrid approach to stance and emotion detection. Few-shot prototypical networks classify stance and affect based on small sets of labeled examples, capturing nuanced positions like punitive loyalists, loyal critics, ironic distancers, and defenders of creative autonomy.

Third, we treat witch-hunt episodes as sequences and seek to identify their internal grammar. Using sequence-rule mining and causal graph techniques, we extract recurrent templates that link trigger events (for example, a controversial historical detail in the game or a developer statement), evidence-selection practices (screenshots, clipped scenes, intertextual references), and moral qualification ("insulting China," "distorting history," "betraying Han people," "traitor within"). This allows us to formalize distinctive narrative units—initial accusation, evidence montage, call to judgment, demand for sanction or boycott—and describe the transition rules that connect them into a reproducible script. Finally, to address RQ3, we examine how users narrate their own participation through self-positioning moves, meta-commentary on "doing justice," and reflections on visibility and rewards, which we relate to basic engagement metrics (views, likes, shares, comments, and where available, monetization signals).

We expect several key findings. First, nationalist witch-hunts around Wuchang: Fallen Feathers will exhibit a templated content logic that is both highly visualizable (via timelines and flow diagrams) and easily replicable. The same sequence—discovery,

accusation, circulation of "evidence," collective moral judgment, and demands for sanction—is likely to recur across platforms and creators with limited variation. Second, accusations framed through Han nationalism and "insulting China" will crystallize into standardized routines. From the moment a potentially problematic element is identified, familiar moves follow: attaching nationalist labels, amplifying worst-case interpretations, and escalating from critique to moral indictment. These routines are optimized for platform visibility, compressing complex historical questions into emotionally charged, shareable packages.

Third, audiences will derive affective rewards from the experience of righteous struggle. Even non-initiating participants obtain emotional payoffs—indignation, catharsis, belonging—through watching, endorsing, and lightly contributing to witch-hunt content. This helps explain why such events can sustain participation across multiple cycles and platforms despite the limited arrival of new information. Fourth, the relationship between performers and audiences will stabilize into recognizable roles that facilitate batch production and cross-platform migration. Influencers and mid-tier creators repeatedly occupy positions as moral entrepreneurs and "case builders," while ordinary users function as jurors, chorus, and amplifiers. As these roles consolidate, nationalism becomes a dependable script for generating engagement and potential income, incentivizing creators to reproduce similar controversies in future cases.

By theorizing cyber-nationalist witch-hunts as both discursive formations and content-production templates, this study contributes to three ongoing debates. For nationalism studies, it highlights how Han nationalism and broader national identity politics are enacted through standardized, platform-optimized repertoires rather than isolated outbursts. For platform and media studies, it shows how recommendation systems and engagement metrics act as semantic infrastructures that privilege specific moralized, othering narratives. For computational social science, it demonstrates how embedding-based clustering, few-shot classification, and sequence mining can be integrated with frame and role analysis to recover the patterned structure of digital campaigns.